

# 1 (XGBoost Classifier with OneHot Encoding)

February 17, 2020

## 0.0.1 Importing Relevant Libraries and Data

```
[14]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

from sklearn.model_selection import train_test_split, GridSearchCV,
    ↳ cross_val_score
from sklearn.preprocessing import OneHotEncoder
from sklearn.metrics import classification_report, confusion_matrix

from xgboost import XGBRFClassifier
```

```
[15]: train_data = pd.read_csv('credit_card_default_train.csv')
test_data = pd.read_csv('credit_card_default_test.csv')
```

```
[16]: cat_cols =
    ↳ ['Gender', 'EDUCATION_STATUS', 'MARITAL_STATUS', 'AGE', 'PAY_JULY', 'PAY_AUG', 'PAY_SEP', 'PAY_OCT']
target = 'NEXT_MONTH_DEFAULT'
ID = 'Client_ID'
num_cols = [col for col in train_data.columns.tolist() if col not in cat_cols
    ↳ + [target] + [ID]]
```

## 0.0.2 Feature Engineering

```
[17]: def Paid_Due_July(row):
    if row['PAID_AMT_JULY'] == 0:
        val = row['DUE_AMT_JULY']
    else:
        val = row['DUE_AMT_JULY']/row['PAID_AMT_JULY']
    return val

def Paid_Due_Aug(row):
```

```

    if row['PAID_AMT_AUG'] == 0:
        val = row['DUE_AMT_AUG']
    else:
        val = row['DUE_AMT_AUG']/row['PAID_AMT_AUG']
    return val

def Paid_Due_Sep(row):
    if row['PAID_AMT_SEP'] == 0:
        val = row['DUE_AMT_SEP']
    else:
        val = row['DUE_AMT_SEP']/row['PAID_AMT_SEP']
    return val

def Paid_Due_Oct(row):
    if row['PAID_AMT_OCT'] == 0:
        val = row['DUE_AMT_OCT']
    else:
        val = row['DUE_AMT_OCT']/row['PAID_AMT_OCT']
    return val

def Paid_Due_Nov(row):
    if row['PAID_AMT_NOV'] == 0:
        val = row['DUE_AMT_NOV']
    else:
        val = row['DUE_AMT_NOV']/row['PAID_AMT_NOV']
    return val

def Paid_Due_Dec(row):
    if row['PAID_AMT_DEC'] == 0:
        val = row['DUE_AMT_DEC']
    else:
        val = row['DUE_AMT_DEC']/row['PAID_AMT_DEC']
    return val

```

```

[18]: train_data['PAID_DUE_JULY'] = train_data.apply(Paid_Due_July, axis=1)
train_data['PAID_DUE_AUG'] = train_data.apply(Paid_Due_Aug, axis=1)
train_data['PAID_DUE_SEP'] = train_data.apply(Paid_Due_Sep, axis=1)
train_data['PAID_DUE_OCT'] = train_data.apply(Paid_Due_Oct, axis=1)
train_data['PAID_DUE_NOV'] = train_data.apply(Paid_Due_Nov, axis=1)
train_data['PAID_DUE_DEC'] = train_data.apply(Paid_Due_Dec, axis=1)

test_data['PAID_DUE_JULY'] = test_data.apply(Paid_Due_July, axis=1)
test_data['PAID_DUE_AUG'] = test_data.apply(Paid_Due_Aug, axis=1)
test_data['PAID_DUE_SEP'] = test_data.apply(Paid_Due_Sep, axis=1)
test_data['PAID_DUE_OCT'] = test_data.apply(Paid_Due_Oct, axis=1)
test_data['PAID_DUE_NOV'] = test_data.apply(Paid_Due_Nov, axis=1)
test_data['PAID_DUE_DEC'] = test_data.apply(Paid_Due_Dec, axis=1)

```

```
[19]: train_data['PAY_TOT'] = train_data['PAY_JULY'] + train_data['PAY_AUG'] +
      ↪train_data['PAY_SEP'] + train_data['PAY_OCT'] + train_data['PAY_NOV'] +
      ↪train_data['PAY_DEC']
      test_data['PAY_TOT'] = test_data['PAY_JULY'] + test_data['PAY_AUG'] +
      ↪test_data['PAY_SEP'] + test_data['PAY_OCT'] + test_data['PAY_NOV'] +
      ↪test_data['PAY_DEC']
```

```
[20]: def isZero(row):
      if row['PAY_TOT'] ==0:
          val = 1
      else:
          val = 0
      return val
```

```
[21]: train_data['PAY_TOT_0'] = train_data.apply(isZero, axis=1)
      test_data['PAY_TOT_0'] = test_data.apply(isZero, axis=1)
```

```
[22]: # creating instance of one-hot-encoder
      enc = OneHotEncoder(handle_unknown='ignore')

      # passing bridge-types-cat column (label encoded values of bridge_types)

      train_df = pd.DataFrame(enc.
      ↪fit_transform(train_data[['Balance_Limit_V1','EDUCATION_STATUS','MARITAL_STATUS','AGE','Gen
      ↪toarray()

      enc = OneHotEncoder(handle_unknown='ignore')

      # passing bridge-types-cat column (label encoded values of bridge_types)

      test_df = pd.DataFrame(enc.
      ↪fit_transform(test_data[['Balance_Limit_V1','EDUCATION_STATUS','MARITAL_STATUS','AGE','Gen
      ↪toarray())
```

```
[23]: train_data = pd.concat([train_data,train_df],axis=1)
      test_data = pd.concat([test_data,test_df],axis=1)
```

### 0.0.3 Model Selection and Evaluation

```
[24]: X = train_data.
      ↪drop([target,ID,'Balance_Limit_V1','Gender','EDUCATION_STATUS','MARITAL_STATUS','AGE'],axis
      y = train_data[target]

      random_grid = {'n_estimators': [100,200,400],
                     'learning_rate': [0.1,0.01,0.03]}
```

```
grid = GridSearchCV(XGBRFClassifier(),random_grid,refit=True,verbose=3,cv=3)
grid.fit(X,y)
```

Fitting 3 folds for each of 9 candidates, totalling 27 fits

[CV] learning\_rate=0.1, n\_estimators=100 ...

[Parallel(n\_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.

[CV] . learning\_rate=0.1, n\_estimators=100, score=0.811, total= 2.5s

[CV] learning\_rate=0.1, n\_estimators=100 ...

[Parallel(n\_jobs=1)]: Done 1 out of 1 | elapsed: 2.4s remaining: 0.0s

[CV] . learning\_rate=0.1, n\_estimators=100, score=0.824, total= 2.3s

[CV] learning\_rate=0.1, n\_estimators=100 ...

[Parallel(n\_jobs=1)]: Done 2 out of 2 | elapsed: 4.7s remaining: 0.0s

[CV] . learning\_rate=0.1, n\_estimators=100, score=0.823, total= 2.3s

[CV] learning\_rate=0.1, n\_estimators=200 ...

[CV] . learning\_rate=0.1, n\_estimators=200, score=0.811, total= 4.3s

[CV] learning\_rate=0.1, n\_estimators=200 ...

[CV] . learning\_rate=0.1, n\_estimators=200, score=0.824, total= 4.4s

[CV] learning\_rate=0.1, n\_estimators=200 ...

[CV] . learning\_rate=0.1, n\_estimators=200, score=0.823, total= 4.4s

[CV] learning\_rate=0.1, n\_estimators=400 ...

[CV] . learning\_rate=0.1, n\_estimators=400, score=0.811, total= 8.6s

[CV] learning\_rate=0.1, n\_estimators=400 ...

[CV] . learning\_rate=0.1, n\_estimators=400, score=0.824, total= 8.5s

[CV] learning\_rate=0.1, n\_estimators=400 ...

[CV] . learning\_rate=0.1, n\_estimators=400, score=0.823, total= 8.7s

[CV] learning\_rate=0.01, n\_estimators=100 ...

[CV] learning\_rate=0.01, n\_estimators=100, score=0.811, total= 2.2s

[CV] learning\_rate=0.01, n\_estimators=100 ...

[CV] learning\_rate=0.01, n\_estimators=100, score=0.824, total= 2.2s

[CV] learning\_rate=0.01, n\_estimators=100 ...

[CV] learning\_rate=0.01, n\_estimators=100, score=0.823, total= 2.2s

[CV] learning\_rate=0.01, n\_estimators=200 ...

[CV] learning\_rate=0.01, n\_estimators=200, score=0.811, total= 4.4s

[CV] learning\_rate=0.01, n\_estimators=200 ...

[CV] learning\_rate=0.01, n\_estimators=200, score=0.824, total= 4.4s

[CV] learning\_rate=0.01, n\_estimators=200 ...

[CV] learning\_rate=0.01, n\_estimators=200, score=0.823, total= 4.4s

[CV] learning\_rate=0.01, n\_estimators=400 ...

[CV] learning\_rate=0.01, n\_estimators=400, score=0.811, total= 8.6s

[CV] learning\_rate=0.01, n\_estimators=400 ...

[CV] learning\_rate=0.01, n\_estimators=400, score=0.824, total= 8.7s

[CV] learning\_rate=0.01, n\_estimators=400 ...

[CV] learning\_rate=0.01, n\_estimators=400, score=0.823, total= 8.5s

[CV] learning\_rate=0.03, n\_estimators=100 ...

```

[CV] learning_rate=0.03, n_estimators=100, score=0.811, total= 2.5s
[CV] learning_rate=0.03, n_estimators=100 ...
[CV] learning_rate=0.03, n_estimators=100, score=0.824, total= 2.5s
[CV] learning_rate=0.03, n_estimators=100 ...
[CV] learning_rate=0.03, n_estimators=100, score=0.823, total= 2.3s
[CV] learning_rate=0.03, n_estimators=200 ...
[CV] learning_rate=0.03, n_estimators=200, score=0.811, total= 4.5s
[CV] learning_rate=0.03, n_estimators=200 ...
[CV] learning_rate=0.03, n_estimators=200, score=0.824, total= 4.5s
[CV] learning_rate=0.03, n_estimators=200 ...
[CV] learning_rate=0.03, n_estimators=200, score=0.823, total= 4.7s
[CV] learning_rate=0.03, n_estimators=400 ...
[CV] learning_rate=0.03, n_estimators=400, score=0.811, total= 8.8s
[CV] learning_rate=0.03, n_estimators=400 ...
[CV] learning_rate=0.03, n_estimators=400, score=0.824, total= 8.7s
[CV] learning_rate=0.03, n_estimators=400 ...
[CV] learning_rate=0.03, n_estimators=400, score=0.823, total= 9.4s

[Parallel(n_jobs=1)]: Done 27 out of 27 | elapsed: 2.3min finished

```

```

[24]: GridSearchCV(cv=3, error_score=nan,
                  estimator=XGBRFClassifier(base_score=0.5, colsample_bylevel=1,
                                           colsample_bynode=0.8, colsample_bytree=1,
                                           gamma=0, learning_rate=1,
                                           max_delta_step=0, max_depth=3,
                                           min_child_weight=1, missing=None,
                                           n_estimators=100, n_jobs=1, nthread=None,
                                           objective='binary:logistic',
                                           random_state=0, reg_alpha=0,
                                           reg_lambda=1, scale_pos_weight=1,
                                           seed=None, silent=None, subsample=0.8,
                                           verbosity=1),
                  iid='deprecated', n_jobs=None,
                  param_grid={'learning_rate': [0.1, 0.01, 0.03],
                              'n_estimators': [100, 200, 400]},
                  pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                  scoring=None, verbose=3)

```

```

[25]: print(grid.best_params_)
      print(grid.best_score_)

```

```

{'learning_rate': 0.1, 'n_estimators': 400}
0.8194166666666667

```

```

[27]: model_xgb = XGBRFClassifier(n_estimators=400, learning_rate=0.1)
      #model_rf.fit(train_data.
      ↪ drop([target, ID, 'Balance_Limit_V1', 'Gender', 'EDUCATION_STATUS', 'MARITAL_STATUS', 'AGE'], axis
      #model_rf.fit(X_train, y_train)

```

```
scores = cross_val_score(model_xgb, X, y, cv=5, scoring='accuracy')
print(scores)
print (np.mean(scores))
```

```
[0.80541667 0.81229167 0.82375      0.83104167 0.821875   ]
0.818875
```

```
[28]: X_train, X_test, y_train, y_test = train_test_split(train_data.
↳drop([target,ID,'Balance_Limit_V1','Gender','EDUCATION_STATUS','MARITAL_STATUS','AGE'],axis
train_data[target],
↳test_size=0.30)
model_xgb.fit(X_train,y_train)
```

```
[28]: XGBRFClassifier(base_score=0.5, colsample_bylevel=1, colsample_bynode=0.8,
colsample_bytree=1, gamma=0, learning_rate=0.1,
max_delta_step=0, max_depth=3, min_child_weight=1, missing=None,
n_estimators=400, n_jobs=1, nthread=None,
objective='binary:logistic', random_state=0, reg_alpha=0,
reg_lambda=1, scale_pos_weight=1, seed=None, silent=None,
subsample=0.8, verbosity=1)
```

```
[29]: preds_xgb = model_xgb.predict(X_test)

print(classification_report(y_test,preds_xgb))
print ('\n')
print(confusion_matrix(y_test,preds_xgb))
```

	precision	recall	f1-score	support
0	0.83	0.96	0.89	5612
1	0.69	0.32	0.44	1588
accuracy			0.82	7200
macro avg	0.76	0.64	0.67	7200
weighted avg	0.80	0.82	0.79	7200

```
[[5385  227]
 [1073  515]]
```

#### 0.0.4 Implementation of the model in test data

```
[33]: model_xgb = XGBRFClassifier(n_estimators=400, learning_rate = 0.1)
      model_xgb.fit(train_data.
      ↪drop([target, ID, 'Balance_Limit_V1', 'Gender', 'EDUCATION_STATUS', 'MARITAL_STATUS', 'AGE'], axis=1))

[33]: XGBRFClassifier(base_score=0.5, colsample_bylevel=1, colsample_bynode=0.8,
      colsample_bytree=1, gamma=0, learning_rate=0.1,
      max_delta_step=0, max_depth=3, min_child_weight=1, missing=None,
      n_estimators=400, n_jobs=1, nthread=None,
      objective='binary:logistic', random_state=0, reg_alpha=0,
      reg_lambda=1, scale_pos_weight=1, seed=None, silent=None,
      subsample=0.8, verbosity=1)

[35]: preds_xgb = model_xgb.predict(test_data.
      ↪drop([ID, 'Balance_Limit_V1', 'Gender', 'EDUCATION_STATUS', 'MARITAL_STATUS', 'AGE'], axis=1))

[36]: sample_submission = pd.DataFrame(columns=[ID, target])
      sample_submission[ID] = test_data[ID]
      sample_submission[target] = preds_xgb

[37]: submission = sample_submission.to_csv('data-storm-day1-3.csv', index = None)
```